

# Capstone\_Project

November 5, 2020

## 1 Capstone Project

### 1.1 Image classifier for the SVHN dataset

#### 1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

#### 1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (you could download the notebook with File -> Download .ipynb, open the notebook locally, and then File -> Download as -> PDF via LaTeX), and then submit this pdf for review.

#### 1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
[ ]: import tensorflow as tf
import pandas as pd
```

For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning". NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

The train and test datasets required for this project can be downloaded from [here](#) and [here](#). Once unzipped, you will have two files: `train_32x32.mat` and `test_32x32.mat`. You should store these files in Drive for use in this Colab notebook.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

```
[ ]: # Downloading the Dataset
```

```
! wget http://ufldl.stanford.edu/housenumbers/train.tar.gz
! wget http://ufldl.stanford.edu/housenumbers/test.tar.gz
```

```
--2020-11-05 17:49:05-- http://ufldl.stanford.edu/housenumbers/train.tar.gz
Resolving ufldl.stanford.edu (ufldl.stanford.edu)... 171.64.68.10
Connecting to ufldl.stanford.edu (ufldl.stanford.edu)|171.64.68.10|:80...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 404141560 (385M) [application/x-gzip]
Saving to: train.tar.gz
```

```
train.tar.gz          100%[=====>] 385.42M  3.50MB/s    in 54s
```

```
2020-11-05 17:49:59 (7.18 MB/s) - train.tar.gz saved [404141560/404141560]
```

```
--2020-11-05 17:49:59-- http://ufldl.stanford.edu/housenumbers/test.tar.gz
Resolving ufldl.stanford.edu (ufldl.stanford.edu)... 171.64.68.10
Connecting to ufldl.stanford.edu (ufldl.stanford.edu)|171.64.68.10|:80...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 276555967 (264M) [application/x-gzip]
Saving to: test.tar.gz
```

```
test.tar.gz           100%[=====>] 263.74M  8.91MB/s    in 1m 53s
```

```
2020-11-05 17:51:52 (2.34 MB/s) - test.tar.gz saved [276555967/276555967]
```

```
[ ]: # Extracting Test Images
```

```
! tar -xzf test.tar.gz
```

```
[ ]: # Extracting Train Images
```

```
! tar -xzf train.tar.gz
```

```
[ ]: ! pip install mat73
```

```
Collecting mat73
```

```
  Downloading https://files.pythonhosted.org/packages/95/f4/175fd094d338c5eb3b33
  8268a301a1109c72f9ed92fa6783cdbc5de40b/mat73-0.45-py3-none-any.whl
```

```
Requirement already satisfied: h5py in /usr/local/lib/python3.6/dist-packages
(from mat73) (2.10.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages
(from mat73) (1.18.5)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages
(from h5py->mat73) (1.15.0)
Installing collected packages: mat73
Successfully installed mat73-0.45
```

```
[ ]: from mat73 import loadmat

[ ]: # Load the dataset from your Drive folder

MATRIX = {}

MATRIX["train"] = loadmat('train/digitStruct.mat')
MATRIX["test"] = loadmat('test/digitStruct.mat')
```

Both train and test are dictionaries with keys X and y for the input images and labels respectively.

## 1.2 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
[ ]: SIZE = {}

SIZE["train"] = len(MATRIX["train"]["digitStruct"]["name"])
SIZE["test"] = len(MATRIX["test"]["digitStruct"]["name"])
```

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
from skimage import transform
from skimage.color import rgb2gray
from collections import Counter
```

```
[ ]: plt.rcParams["figure.figsize"] = 20, 12

for idx in range(20):
    plt.subplot(4, 5, idx+1)
    e = np.random.randint(0, SIZE["train"])
    img = plt.imread("train/{}.png".format(e+1))
    plt.imshow(img)
```

```
plt.axis('off')
plt.title((np.array(MATRIX["train"] ["digitStruct"] ["bbox"] [e] ["label"],
→dtype=np.int16)%10).tolist())
plt.show()

plt.rcParams["figure.figsize"] = 5, 4
```



```
[ ]: """
The function does these operation on the follwoing sequences:
+ For Image:
    + Get the best crop position so that all the digits are included.
    + Convert image into gray scale.
    + Resize the image into (32, 32)
+ For Label:
    + Return Labels as a list
"""

data_point = lambda img, data: (
    transform.resize(image = rgb2gray(img[np.
→abs(data.top).min().astype(np.int16):(np.abs(data.top) + np.abs(data.
→height)).max().astype(np.int16),
    np.
→abs(data.left).min().astype(np.int16):(np.abs(data.left) + np.abs(data.
→width)).max().astype(np.int16)]),
    output_shape=(32, 32)),
    [abs(int(x)%10) for x in data.label]
```

```

        ) \
        if type(data.label) == list else \
        (
            transform.resize(image = rgb2gray(
→img[int(abs(data.top)):int(abs(data.top) + abs(data.height)),
→int(abs(data.left)):int(abs(data.left) + abs(data.width))]
            ),
            output_shape=(32, 32)),
            [abs(int(data.label)%10)]
        )

```

```

[ ]: def fill_arrays(path):
    """
    This function takes the path of the folder where iamges are located.
    It prepares the data so that it can be used for data anlaysis and mdoelling.
    """

    image_array = []
    label_array = []
    i = 0

    while(i < SIZE[path]):
        img = plt.imread("{} / {}".format(path,
→MATRIX[path]["digitStruct"]["name"][i]))
        img, labels = data_point(img, MATRIX[path]["digitStruct"]["bbox"][i])

        image_array.append(img)
        label_array.append(labels)

        i += 1

    return image_array, label_array

```

```

[ ]: train_images, train_labels = fill_arrays("train")
    test_images, test_labels = fill_arrays("test")

```


```

[ ]: plt.rcParams["figure.figsize"] = 20, 12

for idx in range(20):
    plt.subplot(4, 5, idx+1)
    e = np.random.randint(0, SIZE["train"])
    plt.imshow(train_images[e], cmap="binary")
    plt.axis('off')
    plt.title(train_labels[e])
plt.show()

```

```
plt.rcParams["figure.figsize"] = 5, 4
```




```
[ ]: # Frequency of x digit numbers in training set.

train_dist = list(map(lambda x: len(x), train_labels))

plt.hist(train_dist, bins=6, density=True)
plt.show()

Counter(train_dist)
```



```
[ ]: Counter({1: 5137, 2: 18130, 3: 8691, 4: 1434, 5: 9, 6: 1})
```


```
[ ]: # Frequency of x digit numbers in training set.
```

```
test_dist = list(map(lambda x: len(x), test_labels))
```

```
plt.hist(test_dist, bins=6, density=True, )
```

```
plt.show()
```

```
Counter(test_dist)
```



```
[ ]: Counter({1: 2483, 2: 8356, 3: 2081, 4: 146, 5: 2})
```

```
[ ]: # NOTE: The dataset is hugely imbalanced with respect to the number of digits
      → in the images.
```

```
[ ]: from sklearn.preprocessing import MultiLabelBinarizer
```

```
[ ]: X_train = np.array(train_images)
     X_test = np.array(test_images)

     # Create MultiLabelBinarizer object
     one_hot = MultiLabelBinarizer()
     one_hot.fit(train_labels)

     Y_train = one_hot.transform(train_labels)
     Y_test = one_hot.transform(test_labels)
```

### 1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.



- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
[ ]: from tensorflow.keras import layers, models, optimizers, metrics, losses, \
      ↪callbacks
```

```
[ ]: def SequentialModel(input_shape):
      """
      NOTE: Since there are more than one digits present in each image, hence I
      ↪am going to treat this problem as
          multi label classification where an image can have more than one
      ↪labels. For this case, Softmax activation
          doesn't make any sense in the final layer. It is suitable for
      ↪multiclass classification but not for multilabel
          classification. Therefore, I am going to use Sigmoid activation in
      ↪final layer with 10 units which encodes
          the probability of each digit by a Bernoulli variable.
          Same reason goes for the model loss as well. It doesn't make any
      ↪sense to use Categorical cross entropy which is
          suitable for multi-class classification, but not for multi-label
      ↪classification.
      """
      model = models.Sequential([
          layers.InputLayer(input_shape=input_shape),
          layers.Flatten(),
          layers.Dense(60, activation="relu"),
          layers.Dense(100, activation="relu"),
          layers.Dense(100, activation="relu"),
          layers.Dense(10, activation="sigmoid")
      ])

      model.compile(optimizer = "adam",
                    loss = losses.BinaryCrossentropy(),
                    metrics = [metrics.BinaryAccuracy("accuracy"), metrics.
      ↪Precision(name="precision"), metrics.Recall(name="recall")])

      return model
```

```
[ ]: model_1 = SequentialModel(input_shape=(32, 32))

      model_1.summary()
```

Model: "sequential\_11"

Layer (type)	Output Shape	Param #
flatten_11 (Flatten)	(None, 1024)	0
dense_42 (Dense)	(None, 60)	61500
dense_43 (Dense)	(None, 100)	6100
dense_44 (Dense)	(None, 100)	10100
dense_45 (Dense)	(None, 10)	1010

Total params: 78,710  
 Trainable params: 78,710  
 Non-trainable params: 0

```
[ ]: ! rm -rf sequential_model_checkpoint
```

```
[ ]: # Callbacks
```

```

seq_path = "sequential_model_checkpoints/checkpoint"
ckpt_seq = callbacks.ModelCheckpoint(filepath=seq_path,
                                     monitor="val_accuracy",
                                     verbose=True,
                                     mode="max",
                                     save_freq="epoch",
                                     save_weights_only=True,
                                     save_best_only=True)

csvl_seq = callbacks.CSVLogger(filename="sequential_model_CSV.csv")

```

```
[ ]: # Model Training
```

```

history = model_1.fit(x = X_train,
                      y = Y_train,
                      batch_size=300,
                      epochs=50,
                      validation_split = 0.15,
                      callbacks = [ckpt_seq, csvl_seq],
                      verbose = 2)

```

Epoch 1/50

Epoch 00001: val\_accuracy improved from -inf to 0.79563, saving model to sequential\_model\_checkpoints/checkpoint  
 95/95 - 1s - loss: 0.5136 - accuracy: 0.7918 - precision: 0.2496 - recall:

0.0134 - val\_loss: 0.5065 - val\_accuracy: 0.7956 - val\_precision: 0.0000e+00 -  
val\_recall: 0.0000e+00  
Epoch 2/50

Epoch 00002: val\_accuracy did not improve from 0.79563  
95/95 - 0s - loss: 0.4992 - accuracy: 0.7972 - precision: 0.0000e+00 - recall:  
0.0000e+00 - val\_loss: 0.4955 - val\_accuracy: 0.7956 - val\_precision: 0.0000e+00  
- val\_recall: 0.0000e+00  
Epoch 3/50

Epoch 00003: val\_accuracy improved from 0.79563 to 0.79765, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4852 - accuracy: 0.7977 - precision: 0.7317 - recall:  
0.0036 - val\_loss: 0.4768 - val\_accuracy: 0.7976 - val\_precision: 0.7563 -  
val\_recall: 0.0145  
Epoch 4/50

Epoch 00004: val\_accuracy improved from 0.79765 to 0.80347, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4655 - accuracy: 0.8034 - precision: 0.6883 - recall:  
0.0561 - val\_loss: 0.4605 - val\_accuracy: 0.8035 - val\_precision: 0.6546 -  
val\_recall: 0.0812  
Epoch 5/50

Epoch 00005: val\_accuracy improved from 0.80347 to 0.80868, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4538 - accuracy: 0.8086 - precision: 0.6706 - recall:  
0.1099 - val\_loss: 0.4502 - val\_accuracy: 0.8087 - val\_precision: 0.6972 -  
val\_recall: 0.1129  
Epoch 6/50

Epoch 00006: val\_accuracy improved from 0.80868 to 0.81371, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4400 - accuracy: 0.8142 - precision: 0.6970 - recall:  
0.1480 - val\_loss: 0.4395 - val\_accuracy: 0.8137 - val\_precision: 0.6933 -  
val\_recall: 0.1587  
Epoch 7/50

Epoch 00007: val\_accuracy improved from 0.81371 to 0.81642, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4290 - accuracy: 0.8187 - precision: 0.7048 - recall:  
0.1820 - val\_loss: 0.4339 - val\_accuracy: 0.8164 - val\_precision: 0.7056 -  
val\_recall: 0.1746  
Epoch 8/50

Epoch 00008: val\_accuracy improved from 0.81642 to 0.82032, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4196 - accuracy: 0.8223 - precision: 0.7077 - recall:

0.2111 - val\_loss: 0.4236 - val\_accuracy: 0.8203 - val\_precision: 0.7132 -  
val\_recall: 0.2020  
Epoch 9/50

Epoch 00009: val\_accuracy did not improve from 0.82032  
95/95 - 0s - loss: 0.4122 - accuracy: 0.8255 - precision: 0.7103 - recall:  
0.2352 - val\_loss: 0.4201 - val\_accuracy: 0.8194 - val\_precision: 0.6414 -  
val\_recall: 0.2640  
Epoch 10/50

Epoch 00010: val\_accuracy improved from 0.82032 to 0.82548, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4062 - accuracy: 0.8278 - precision: 0.7126 - recall:  
0.2528 - val\_loss: 0.4120 - val\_accuracy: 0.8255 - val\_precision: 0.6705 -  
val\_recall: 0.2873  
Epoch 11/50

Epoch 00011: val\_accuracy improved from 0.82548 to 0.82766, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.4024 - accuracy: 0.8297 - precision: 0.7169 - recall:  
0.2648 - val\_loss: 0.4099 - val\_accuracy: 0.8277 - val\_precision: 0.6712 -  
val\_recall: 0.3072  
Epoch 12/50

Epoch 00012: val\_accuracy improved from 0.82766 to 0.82834, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3980 - accuracy: 0.8314 - precision: 0.7153 - recall:  
0.2797 - val\_loss: 0.4055 - val\_accuracy: 0.8283 - val\_precision: 0.7001 -  
val\_recall: 0.2800  
Epoch 13/50

Epoch 00013: val\_accuracy improved from 0.82834 to 0.82930, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3942 - accuracy: 0.8331 - precision: 0.7205 - recall:  
0.2890 - val\_loss: 0.4005 - val\_accuracy: 0.8293 - val\_precision: 0.7163 -  
val\_recall: 0.2727  
Epoch 14/50

Epoch 00014: val\_accuracy improved from 0.82930 to 0.83331, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3917 - accuracy: 0.8341 - precision: 0.7182 - recall:  
0.2992 - val\_loss: 0.3941 - val\_accuracy: 0.8333 - val\_precision: 0.7453 -  
val\_recall: 0.2801  
Epoch 15/50

Epoch 00015: val\_accuracy did not improve from 0.83331  
95/95 - 0s - loss: 0.3900 - accuracy: 0.8348 - precision: 0.7206 - recall:  
0.3024 - val\_loss: 0.4027 - val\_accuracy: 0.8278 - val\_precision: 0.6751 -

val\_recall: 0.3033  
Epoch 16/50

Epoch 00016: val\_accuracy improved from 0.83331 to 0.83349, saving model to sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3861 - accuracy: 0.8370 - precision: 0.7273 - recall: 0.3136 - val\_loss: 0.3902 - val\_accuracy: 0.8335 - val\_precision: 0.7273 - val\_recall: 0.2964  
Epoch 17/50

Epoch 00017: val\_accuracy improved from 0.83349 to 0.83365, saving model to sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3844 - accuracy: 0.8374 - precision: 0.7266 - recall: 0.3174 - val\_loss: 0.3906 - val\_accuracy: 0.8336 - val\_precision: 0.7290 - val\_recall: 0.2961  
Epoch 18/50

Epoch 00018: val\_accuracy improved from 0.83365 to 0.83720, saving model to sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3805 - accuracy: 0.8390 - precision: 0.7314 - recall: 0.3257 - val\_loss: 0.3861 - val\_accuracy: 0.8372 - val\_precision: 0.7161 - val\_recall: 0.3370  
Epoch 19/50

Epoch 00019: val\_accuracy did not improve from 0.83720  
95/95 - 0s - loss: 0.3779 - accuracy: 0.8406 - precision: 0.7368 - recall: 0.3325 - val\_loss: 0.3884 - val\_accuracy: 0.8362 - val\_precision: 0.7080 - val\_recall: 0.3376  
Epoch 20/50

Epoch 00020: val\_accuracy improved from 0.83720 to 0.83732, saving model to sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3769 - accuracy: 0.8410 - precision: 0.7342 - recall: 0.3382 - val\_loss: 0.3890 - val\_accuracy: 0.8373 - val\_precision: 0.7247 - val\_recall: 0.3290  
Epoch 21/50

Epoch 00021: val\_accuracy did not improve from 0.83732  
95/95 - 0s - loss: 0.3733 - accuracy: 0.8424 - precision: 0.7383 - recall: 0.3452 - val\_loss: 0.3869 - val\_accuracy: 0.8359 - val\_precision: 0.7256 - val\_recall: 0.3169  
Epoch 22/50

Epoch 00022: val\_accuracy improved from 0.83732 to 0.84129, saving model to sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3728 - accuracy: 0.8430 - precision: 0.7392 - recall: 0.3488 - val\_loss: 0.3792 - val\_accuracy: 0.8413 - val\_precision: 0.7305 - val\_recall: 0.3541

Epoch 23/50

Epoch 00023: val\_accuracy did not improve from 0.84129

95/95 - 0s - loss: 0.3719 - accuracy: 0.8434 - precision: 0.7387 - recall:  
0.3525 - val\_loss: 0.3791 - val\_accuracy: 0.8402 - val\_precision: 0.7325 -  
val\_recall: 0.3431

Epoch 24/50

Epoch 00024: val\_accuracy improved from 0.84129 to 0.84139, saving model to  
sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3685 - accuracy: 0.8450 - precision: 0.7450 - recall:  
0.3583 - val\_loss: 0.3754 - val\_accuracy: 0.8414 - val\_precision: 0.7275 -  
val\_recall: 0.3580

Epoch 25/50

Epoch 00025: val\_accuracy improved from 0.84139 to 0.84281, saving model to  
sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3666 - accuracy: 0.8461 - precision: 0.7448 - recall:  
0.3666 - val\_loss: 0.3754 - val\_accuracy: 0.8428 - val\_precision: 0.7361 -  
val\_recall: 0.3598

Epoch 26/50

Epoch 00026: val\_accuracy did not improve from 0.84281

95/95 - 0s - loss: 0.3642 - accuracy: 0.8471 - precision: 0.7485 - recall:  
0.3706 - val\_loss: 0.3741 - val\_accuracy: 0.8417 - val\_precision: 0.7257 -  
val\_recall: 0.3628

Epoch 27/50

Epoch 00027: val\_accuracy improved from 0.84281 to 0.84346, saving model to  
sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3623 - accuracy: 0.8484 - precision: 0.7526 - recall:  
0.3759 - val\_loss: 0.3716 - val\_accuracy: 0.8435 - val\_precision: 0.7379 -  
val\_recall: 0.3630

Epoch 28/50

Epoch 00028: val\_accuracy improved from 0.84346 to 0.84502, saving model to  
sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3599 - accuracy: 0.8488 - precision: 0.7505 - recall:  
0.3807 - val\_loss: 0.3689 - val\_accuracy: 0.8450 - val\_precision: 0.7403 -  
val\_recall: 0.3722

Epoch 29/50

Epoch 00029: val\_accuracy did not improve from 0.84502

95/95 - 0s - loss: 0.3576 - accuracy: 0.8501 - precision: 0.7547 - recall:  
0.3861 - val\_loss: 0.3705 - val\_accuracy: 0.8438 - val\_precision: 0.7186 -  
val\_recall: 0.3877

Epoch 30/50

Epoch 00030: val\_accuracy improved from 0.84502 to 0.84542, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3552 - accuracy: 0.8516 - precision: 0.7599 - recall: 0.3922 - val\_loss: 0.3683 - val\_accuracy: 0.8454 - val\_precision: 0.7455 - val\_recall: 0.3699

Epoch 31/50

Epoch 00031: val\_accuracy improved from 0.84542 to 0.84646, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3550 - accuracy: 0.8514 - precision: 0.7550 - recall: 0.3953 - val\_loss: 0.3658 - val\_accuracy: 0.8465 - val\_precision: 0.7473 - val\_recall: 0.3757

Epoch 32/50

Epoch 00032: val\_accuracy did not improve from 0.84646

95/95 - 0s - loss: 0.3523 - accuracy: 0.8522 - precision: 0.7578 - recall: 0.3987 - val\_loss: 0.3680 - val\_accuracy: 0.8454 - val\_precision: 0.7343 - val\_recall: 0.3813

Epoch 33/50

Epoch 00033: val\_accuracy did not improve from 0.84646

95/95 - 0s - loss: 0.3509 - accuracy: 0.8526 - precision: 0.7566 - recall: 0.4025 - val\_loss: 0.3721 - val\_accuracy: 0.8432 - val\_precision: 0.7218 - val\_recall: 0.3788

Epoch 34/50

Epoch 00034: val\_accuracy improved from 0.84646 to 0.84704, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3501 - accuracy: 0.8532 - precision: 0.7587 - recall: 0.4051 - val\_loss: 0.3640 - val\_accuracy: 0.8470 - val\_precision: 0.7342 - val\_recall: 0.3943

Epoch 35/50

Epoch 00035: val\_accuracy improved from 0.84704 to 0.84851, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3476 - accuracy: 0.8546 - precision: 0.7627 - recall: 0.4105 - val\_loss: 0.3611 - val\_accuracy: 0.8485 - val\_precision: 0.7332 - val\_recall: 0.4068

Epoch 36/50

Epoch 00036: val\_accuracy improved from 0.84851 to 0.84857, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3471 - accuracy: 0.8547 - precision: 0.7615 - recall: 0.4127 - val\_loss: 0.3618 - val\_accuracy: 0.8486 - val\_precision: 0.7314 - val\_recall: 0.4094

Epoch 37/50

Epoch 00037: val\_accuracy did not improve from 0.84857

95/95 - 0s - loss: 0.3441 - accuracy: 0.8560 - precision: 0.7651 - recall: 0.4180 - val\_loss: 0.3614 - val\_accuracy: 0.8473 - val\_precision: 0.7533 - val\_recall: 0.3760

Epoch 38/50

Epoch 00038: val\_accuracy improved from 0.84857 to 0.84951, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3424 - accuracy: 0.8565 - precision: 0.7641 - recall: 0.4227 - val\_loss: 0.3579 - val\_accuracy: 0.8495 - val\_precision: 0.7388 - val\_recall: 0.4079

Epoch 39/50

Epoch 00039: val\_accuracy did not improve from 0.84951

95/95 - 0s - loss: 0.3421 - accuracy: 0.8567 - precision: 0.7642 - recall: 0.4243 - val\_loss: 0.3640 - val\_accuracy: 0.8479 - val\_precision: 0.7422 - val\_recall: 0.3919

Epoch 40/50

Epoch 00040: val\_accuracy did not improve from 0.84951

95/95 - 0s - loss: 0.3410 - accuracy: 0.8572 - precision: 0.7648 - recall: 0.4273 - val\_loss: 0.3602 - val\_accuracy: 0.8484 - val\_precision: 0.7185 - val\_recall: 0.4247

Epoch 41/50

Epoch 00041: val\_accuracy improved from 0.84951 to 0.85049, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3410 - accuracy: 0.8575 - precision: 0.7651 - recall: 0.4290 - val\_loss: 0.3569 - val\_accuracy: 0.8505 - val\_precision: 0.7393 - val\_recall: 0.4147

Epoch 42/50

Epoch 00042: val\_accuracy did not improve from 0.85049

95/95 - 0s - loss: 0.3381 - accuracy: 0.8583 - precision: 0.7672 - recall: 0.4325 - val\_loss: 0.3602 - val\_accuracy: 0.8495 - val\_precision: 0.7272 - val\_recall: 0.4218

Epoch 43/50

Epoch 00043: val\_accuracy improved from 0.85049 to 0.85121, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3366 - accuracy: 0.8590 - precision: 0.7679 - recall: 0.4364 - val\_loss: 0.3550 - val\_accuracy: 0.8512 - val\_precision: 0.7317 - val\_recall: 0.4294

Epoch 44/50

Epoch 00044: val\_accuracy improved from 0.85121 to 0.85157, saving model to sequential\_model\_checkpoints/checkpoint

95/95 - 0s - loss: 0.3371 - accuracy: 0.8589 - precision: 0.7675 - recall: 0.4367 - val\_loss: 0.3559 - val\_accuracy: 0.8516 - val\_precision: 0.7415 -



val\_recall: 0.4202  
Epoch 45/50

Epoch 00045: val\_accuracy did not improve from 0.85157  
95/95 - 0s - loss: 0.3353 - accuracy: 0.8601 - precision: 0.7716 - recall:  
0.4403 - val\_loss: 0.3556 - val\_accuracy: 0.8510 - val\_precision: 0.7384 -  
val\_recall: 0.4196  
Epoch 46/50

Epoch 00046: val\_accuracy improved from 0.85157 to 0.85342, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3343 - accuracy: 0.8603 - precision: 0.7699 - recall:  
0.4436 - val\_loss: 0.3504 - val\_accuracy: 0.8534 - val\_precision: 0.7452 -  
val\_recall: 0.4297  
Epoch 47/50

Epoch 00047: val\_accuracy did not improve from 0.85342  
95/95 - 0s - loss: 0.3328 - accuracy: 0.8607 - precision: 0.7719 - recall:  
0.4445 - val\_loss: 0.3529 - val\_accuracy: 0.8531 - val\_precision: 0.7391 -  
val\_recall: 0.4349  
Epoch 48/50

Epoch 00048: val\_accuracy did not improve from 0.85342  
95/95 - 0s - loss: 0.3308 - accuracy: 0.8620 - precision: 0.7744 - recall:  
0.4512 - val\_loss: 0.3562 - val\_accuracy: 0.8506 - val\_precision: 0.7289 -  
val\_recall: 0.4287  
Epoch 49/50

Epoch 00049: val\_accuracy improved from 0.85342 to 0.85378, saving model to  
sequential\_model\_checkpoints/checkpoint  
95/95 - 0s - loss: 0.3316 - accuracy: 0.8611 - precision: 0.7693 - recall:  
0.4503 - val\_loss: 0.3500 - val\_accuracy: 0.8538 - val\_precision: 0.7474 -  
val\_recall: 0.4298  
Epoch 50/50

Epoch 00050: val\_accuracy did not improve from 0.85378  
95/95 - 0s - loss: 0.3295 - accuracy: 0.8627 - precision: 0.7762 - recall:  
0.4535 - val\_loss: 0.3516 - val\_accuracy: 0.8526 - val\_precision: 0.7522 -  
val\_recall: 0.4159

```
[ ]: # Results Visualization
```

```
plt.rcParams["figure.figsize"] = 20, 8

plt.subplot(2, 2, 1)
plt.plot(history.history["loss"], label="Train Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
```

```

plt.legend()
plt.title("Loss Curves")

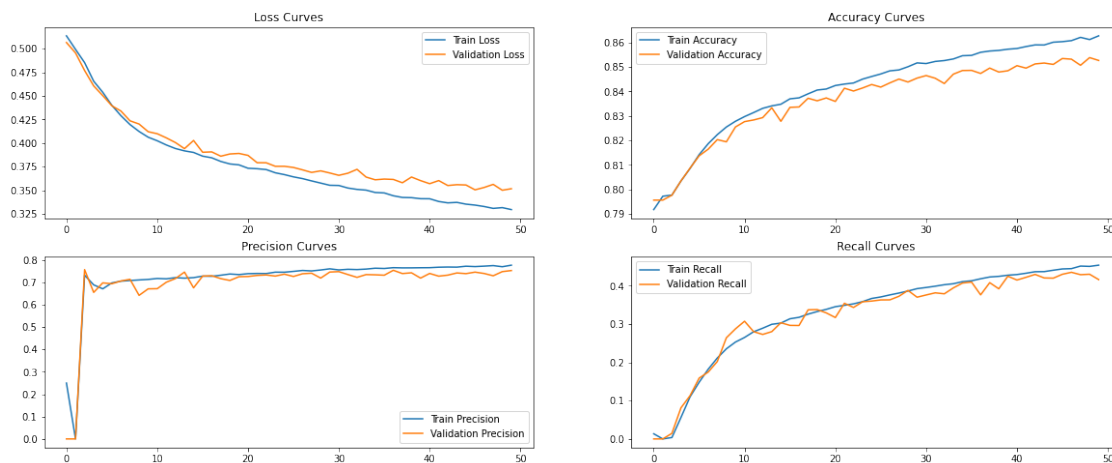
plt.subplot(2, 2, 2)
plt.plot(history.history["accuracy"], label="Train Accuracy")
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.legend()
plt.title("Accuracy Curves")

plt.subplot(2, 2, 3)
plt.plot(history.history["precision"], label="Train Precision")
plt.plot(history.history["val_precision"], label="Validation Precision")
plt.legend()
plt.title("Precision Curves")

plt.subplot(2, 2, 4)
plt.plot(history.history["recall"], label="Train Recall")
plt.plot(history.history["val_recall"], label="Validation Recall")
plt.legend()
plt.title("Recall Curves")

plt.show()

```



```

[ ]: # Evalaution on Test Set

seq_res = model_1.evaluate(x = X_test,
                           y = Y_test,
                           return_dict=True,
                           verbose=False)

```

```

for key in seq_res.keys():
    if key != "loss":
        print("{}: {:.6f}%".format(key.capitalize(), 100*seq_res[key]))
    else:
        print("{}: {:.6f}".format(key.capitalize(), seq_res[key]))

```

Loss: 0.366178  
 Accuracy: 85.306132%  
 Precision: 68.388259%  
 Recall: 40.295342%

## 1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```

[ ]: def CNNModel(input_shape):
    """
    NOTE: Since there are more than one digits present in each image, hence I
    → am going to treat this problem as
        multi label classification where an image can have more than one
    → labels. For this case, Softmax activation
        doesn't make any sense in the final layer. It is suitable for
    → multiclass classification but not for multilabel
        classification. Therefore, I am going to use Sigmoid activation in
    → final layer with 10 units which encodes
        the probability of each digit by a Bernoulli variable.
        Same reason goes for the model loss as well. It doesn't make any
    → sense to use Categorical cross entropy which is
        suitable for multi-class classification, but not for multi-label
    → classification.
    """
    model = models.Sequential([
        layers.InputLayer(input_shape=input_shape),

```

```

        layers.Conv2D(32, kernel_size=(3, 3),
→activation="relu"),
        layers.Dropout(rate=0.2),
        layers.BatchNormalization(),
        layers.MaxPool2D(pool_size=(2,2)),
        layers.Conv2D(64, kernel_size=(3, 3),
→activation="relu"),
        layers.Dropout(rate=0.2),
        layers.BatchNormalization(),
        layers.MaxPool2D(pool_size=(2,2)),
        layers.Flatten(),
        layers.Dense(20, activation="relu"),
        layers.Dense(10, activation="sigmoid")
    ])

    model.compile(optimizer = "adam",
                  loss = losses.BinaryCrossentropy(),
                  metrics = [metrics.BinaryAccuracy("accuracy"), metrics.
→Precision(name="precision"), metrics.Recall(name="recall")])

    return model

```

```

[ ]: model_2 = CNNModel(input_shape=(32, 32, 1))

model_2.summary()

```

Model: "sequential\_13"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 30, 30, 32)	320
dropout_4 (Dropout)	(None, 30, 30, 32)	0
batch_normalization_4 (Batch Normalization)	(None, 30, 30, 32)	128
max_pooling2d_4 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_5 (Conv2D)	(None, 13, 13, 64)	18496
dropout_5 (Dropout)	(None, 13, 13, 64)	0
batch_normalization_5 (Batch Normalization)	(None, 13, 13, 64)	256
max_pooling2d_5 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_13 (Flatten)	(None, 2304)	0

dense_48 (Dense)	(None, 20)	46100
-----		
dense_49 (Dense)	(None, 10)	210
=====		
Total params: 65,510		
Trainable params: 65,318		
Non-trainable params: 192		
-----		

```
[ ]: # Callbacks

cnn_path = "cnn_model_checkpoints/checkpoint"
ckpt_cnn = callbacks.ModelCheckpoint(filepath=cnn_path,
                                     monitor="val_accuracy",
                                     verbose=True,
                                     mode="max",
                                     save_freq="epoch",
                                     save_weights_only=True,
                                     save_best_only=True)

csvl_cnn = callbacks.CSVLogger(filename="cnn_model_CSV.csv")

estop_cnn = callbacks.EarlyStopping(monitor='val_accuracy',
                                    patience=5,
                                    verbose=0,
                                    mode='max',
                                    restore_best_weights=False)

[ ]: history = model_2.fit(x = X_train[..., np.newaxis],
                          y = Y_train,
                          batch_size=100,
                          epochs=50,
                          validation_split = 0.15,
                          callbacks = [ckpt_cnn, csvl_cnn, estop_cnn],
                          verbose = 2)
```

Epoch 1/50

Epoch 00001: val\_accuracy improved from -inf to 0.63055, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.4372 - accuracy: 0.8119 - precision: 0.5869 - recall:  
0.2439 - val\_loss: 0.8093 - val\_accuracy: 0.6306 - val\_precision: 0.2531 -  
val\_recall: 0.4141  
Epoch 2/50

Epoch 00002: val\_accuracy improved from 0.63055 to 0.84568, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.3449 - accuracy: 0.8547 - precision: 0.7319 - recall:

0.4472 - val\_loss: 0.4141 - val\_accuracy: 0.8457 - val\_precision: 0.6647 -  
val\_recall: 0.4943  
Epoch 3/50

Epoch 00003: val\_accuracy improved from 0.84568 to 0.86961, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.3052 - accuracy: 0.8722 - precision: 0.7681 - recall:  
0.5295 - val\_loss: 0.3067 - val\_accuracy: 0.8696 - val\_precision: 0.7944 -  
val\_recall: 0.4883  
Epoch 4/50

Epoch 00004: val\_accuracy did not improve from 0.86961  
284/284 - 2s - loss: 0.2806 - accuracy: 0.8841 - precision: 0.7885 - recall:  
0.5855 - val\_loss: 0.3400 - val\_accuracy: 0.8648 - val\_precision: 0.7493 -  
val\_recall: 0.5086  
Epoch 5/50

Epoch 00005: val\_accuracy improved from 0.86961 to 0.87909, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.2654 - accuracy: 0.8907 - precision: 0.7982 - recall:  
0.6170 - val\_loss: 0.3064 - val\_accuracy: 0.8791 - val\_precision: 0.8212 -  
val\_recall: 0.5220  
Epoch 6/50

Epoch 00006: val\_accuracy improved from 0.87909 to 0.88092, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.2521 - accuracy: 0.8968 - precision: 0.8097 - recall:  
0.6421 - val\_loss: 0.2896 - val\_accuracy: 0.8809 - val\_precision: 0.7640 -  
val\_recall: 0.6039  
Epoch 7/50

Epoch 00007: val\_accuracy improved from 0.88092 to 0.88711, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.2425 - accuracy: 0.9009 - precision: 0.8161 - recall:  
0.6599 - val\_loss: 0.2743 - val\_accuracy: 0.8871 - val\_precision: 0.7912 -  
val\_recall: 0.6081  
Epoch 8/50

Epoch 00008: val\_accuracy did not improve from 0.88711  
284/284 - 2s - loss: 0.2351 - accuracy: 0.9045 - precision: 0.8235 - recall:  
0.6734 - val\_loss: 0.2940 - val\_accuracy: 0.8807 - val\_precision: 0.7732 -  
val\_recall: 0.5891  
Epoch 9/50

Epoch 00009: val\_accuracy improved from 0.88711 to 0.89026, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.2276 - accuracy: 0.9080 - precision: 0.8288 - recall:  
0.6884 - val\_loss: 0.2734 - val\_accuracy: 0.8903 - val\_precision: 0.8189 -

val\_recall: 0.5946  
Epoch 10/50

Epoch 00010: val\_accuracy improved from 0.89026 to 0.89174, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.2201 - accuracy: 0.9113 - precision: 0.8348 - recall:  
0.7011 - val\_loss: 0.2653 - val\_accuracy: 0.8917 - val\_precision: 0.7523 -  
val\_recall: 0.7012  
Epoch 11/50

Epoch 00011: val\_accuracy improved from 0.89174 to 0.89677, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.2134 - accuracy: 0.9148 - precision: 0.8425 - recall:  
0.7135 - val\_loss: 0.2585 - val\_accuracy: 0.8968 - val\_precision: 0.8228 -  
val\_recall: 0.6307  
Epoch 12/50

Epoch 00012: val\_accuracy improved from 0.89677 to 0.90273, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.2078 - accuracy: 0.9164 - precision: 0.8443 - recall:  
0.7207 - val\_loss: 0.2434 - val\_accuracy: 0.9027 - val\_precision: 0.8076 -  
val\_recall: 0.6879  
Epoch 13/50

Epoch 00013: val\_accuracy did not improve from 0.90273  
284/284 - 2s - loss: 0.2019 - accuracy: 0.9190 - precision: 0.8493 - recall:  
0.7302 - val\_loss: 0.2571 - val\_accuracy: 0.8996 - val\_precision: 0.8306 -  
val\_recall: 0.6392  
Epoch 14/50

Epoch 00014: val\_accuracy improved from 0.90273 to 0.90573, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.1981 - accuracy: 0.9208 - precision: 0.8526 - recall:  
0.7370 - val\_loss: 0.2338 - val\_accuracy: 0.9057 - val\_precision: 0.8199 -  
val\_recall: 0.6904  
Epoch 15/50

Epoch 00015: val\_accuracy did not improve from 0.90573  
284/284 - 2s - loss: 0.1929 - accuracy: 0.9234 - precision: 0.8575 - recall:  
0.7461 - val\_loss: 0.2526 - val\_accuracy: 0.9021 - val\_precision: 0.8276 -  
val\_recall: 0.6580  
Epoch 16/50

Epoch 00016: val\_accuracy improved from 0.90573 to 0.90854, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.1886 - accuracy: 0.9253 - precision: 0.8616 - recall:  
0.7523 - val\_loss: 0.2285 - val\_accuracy: 0.9085 - val\_precision: 0.8377 -  
val\_recall: 0.6853

Epoch 17/50

Epoch 00017: val\_accuracy did not improve from 0.90854  
284/284 - 2s - loss: 0.1848 - accuracy: 0.9264 - precision: 0.8632 - recall:  
0.7568 - val\_loss: 0.2432 - val\_accuracy: 0.9055 - val\_precision: 0.8224 -  
val\_recall: 0.6855

Epoch 18/50

Epoch 00018: val\_accuracy improved from 0.90854 to 0.91287, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.1806 - accuracy: 0.9280 - precision: 0.8657 - recall:  
0.7633 - val\_loss: 0.2261 - val\_accuracy: 0.9129 - val\_precision: 0.8524 -  
val\_recall: 0.6938

Epoch 19/50

Epoch 00019: val\_accuracy did not improve from 0.91287  
284/284 - 2s - loss: 0.1777 - accuracy: 0.9300 - precision: 0.8696 - recall:  
0.7704 - val\_loss: 0.2274 - val\_accuracy: 0.9101 - val\_precision: 0.8352 -  
val\_recall: 0.6981

Epoch 20/50

Epoch 00020: val\_accuracy did not improve from 0.91287  
284/284 - 2s - loss: 0.1760 - accuracy: 0.9305 - precision: 0.8717 - recall:  
0.7710 - val\_loss: 0.2223 - val\_accuracy: 0.9128 - val\_precision: 0.8452 -  
val\_recall: 0.7021

Epoch 21/50

Epoch 00021: val\_accuracy improved from 0.91287 to 0.91381, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.1723 - accuracy: 0.9319 - precision: 0.8743 - recall:  
0.7760 - val\_loss: 0.2193 - val\_accuracy: 0.9138 - val\_precision: 0.8386 -  
val\_recall: 0.7161

Epoch 22/50

Epoch 00022: val\_accuracy improved from 0.91381 to 0.91527, saving model to  
cnn\_model\_checkpoints/checkpoint  
284/284 - 2s - loss: 0.1695 - accuracy: 0.9331 - precision: 0.8756 - recall:  
0.7813 - val\_loss: 0.2191 - val\_accuracy: 0.9153 - val\_precision: 0.8589 -  
val\_recall: 0.7005

Epoch 23/50

Epoch 00023: val\_accuracy did not improve from 0.91527  
284/284 - 2s - loss: 0.1673 - accuracy: 0.9340 - precision: 0.8778 - recall:  
0.7836 - val\_loss: 0.2278 - val\_accuracy: 0.9095 - val\_precision: 0.8328 -  
val\_recall: 0.6972

Epoch 24/50

Epoch 00024: val\_accuracy improved from 0.91527 to 0.91766, saving model to



```
cnn_model_checkpoints/checkpoint
284/284 - 2s - loss: 0.1646 - accuracy: 0.9348 - precision: 0.8799 - recall:
0.7859 - val_loss: 0.2113 - val_accuracy: 0.9177 - val_precision: 0.8438 -
val_recall: 0.7327
Epoch 25/50
```

```
Epoch 00025: val_accuracy did not improve from 0.91766
284/284 - 2s - loss: 0.1619 - accuracy: 0.9361 - precision: 0.8824 - recall:
0.7900 - val_loss: 0.2107 - val_accuracy: 0.9168 - val_precision: 0.8415 -
val_recall: 0.7308
Epoch 26/50
```

```
Epoch 00026: val_accuracy did not improve from 0.91766
284/284 - 2s - loss: 0.1598 - accuracy: 0.9372 - precision: 0.8835 - recall:
0.7950 - val_loss: 0.2240 - val_accuracy: 0.9138 - val_precision: 0.8650 -
val_recall: 0.6851
Epoch 27/50
```

```
Epoch 00027: val_accuracy did not improve from 0.91766
284/284 - 2s - loss: 0.1577 - accuracy: 0.9385 - precision: 0.8867 - recall:
0.7990 - val_loss: 0.2333 - val_accuracy: 0.9112 - val_precision: 0.8462 -
val_recall: 0.6909
Epoch 28/50
```

```
Epoch 00028: val_accuracy did not improve from 0.91766
284/284 - 2s - loss: 0.1562 - accuracy: 0.9384 - precision: 0.8863 - recall:
0.7987 - val_loss: 0.2220 - val_accuracy: 0.9144 - val_precision: 0.8545 -
val_recall: 0.7007
Epoch 29/50
```

```
Epoch 00029: val_accuracy did not improve from 0.91766
284/284 - 2s - loss: 0.1535 - accuracy: 0.9401 - precision: 0.8899 - recall:
0.8044 - val_loss: 0.2237 - val_accuracy: 0.9126 - val_precision: 0.8379 -
val_recall: 0.7094
```

```
[ ]: # Reults Visualization

plt.rcParams["figure.figsize"] = 20, 8

plt.subplot(2, 2, 1)
plt.plot(history.history["loss"], label="Train Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.legend()
plt.title("Loss Curves")

plt.subplot(2, 2, 2)
plt.plot(history.history["accuracy"], label="Train Accuracy")
```

```

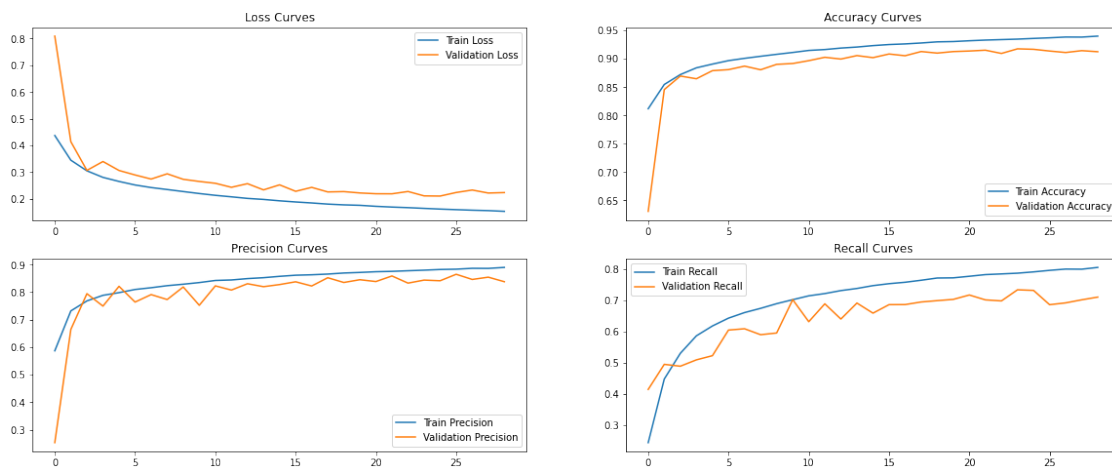
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.legend()
plt.title("Accuracy Curves")

plt.subplot(2, 2, 3)
plt.plot(history.history["precision"], label="Train Precision")
plt.plot(history.history["val_precision"], label="Validation Precision")
plt.legend()
plt.title("Precision Curves")

plt.subplot(2, 2, 4)
plt.plot(history.history["recall"], label="Train Recall")
plt.plot(history.history["val_recall"], label="Validation Recall")
plt.legend()
plt.title("Recall Curves")

plt.show()

```



```
[ ]: # Evalaution on Test Set
```

```

cnn_res = model_2.evaluate(x = X_test[..., np.newaxis],
                           y = Y_test,
                           return_dict=True,
                           verbose=False)

for key in cnn_res.keys():
    if key != "loss":
        print("{}: {:.0.6f}%".format(key.capitalize(), 100*cnn_res[key]))
    else:

```

```
print("{}: {:.6f}".format(key.capitalize(), cnn_res[key]))
```

Loss: 0.230679

Accuracy: 91.500604%

Precision: 81.986928%

Recall: 70.090562%

## 1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

```
[ ]: # Loading best Sequential Model
```

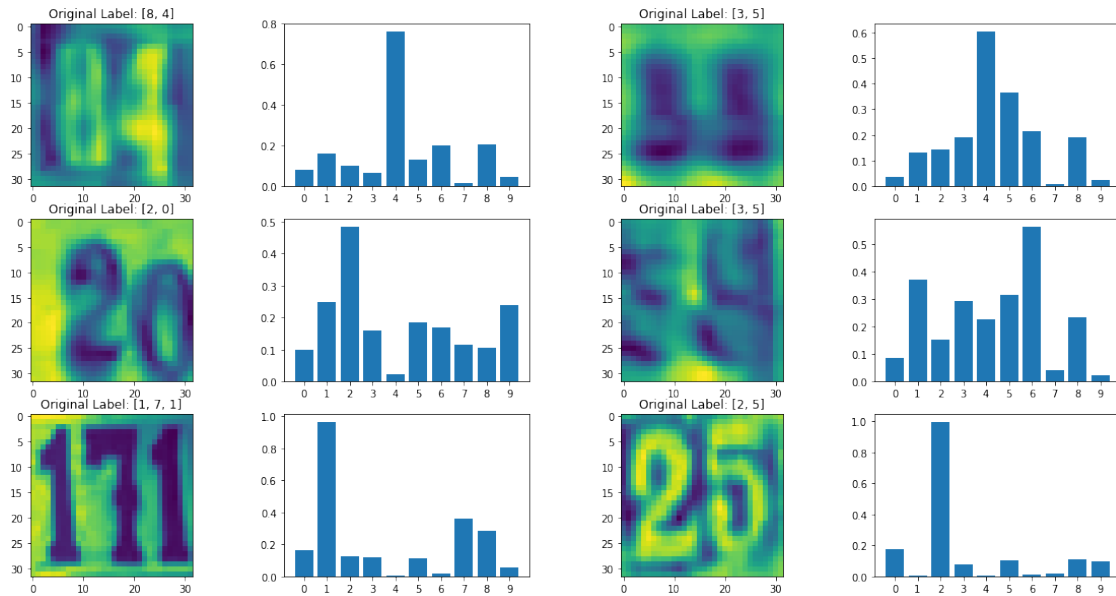
```
best_seq = SequentialModel(input_shape=(32, 32))  
best_seq.load_weights(seq_path)
```

```
[ ]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at  
0x7fc2a80d7048>
```

```
[ ]: # NOTE: I have done this exercise for MultiLabel Classification.  
# Hence, the reviewer should see probability across multiple digits.
```

```
[ ]: # Predictions for best Sequential Model
```

```
plt.rcParams["figure.figsize"] = 20, 10  
  
idx = np.random.randint(0, SIZE["train"], 6)  
  
for i in range(6):  
    plt.subplot(3, 4, 2*i+1)  
    plt.imshow(train_images[idx[i]])  
    plt.title("Original Label: {}".format(train_labels[idx[i]]))  
  
    plt.subplot(3, 4, 2*i + 2)  
    plt.bar(x = np.arange(10), height=best_seq.predict(train_images[idx[i]])[np.  
→newaxis, ...])[0])  
    plt.xticks(np.arange(10))  
  
plt.show()
```



```
[ ]: # Loading best CNN Model
```

```
best_cnn = CNNModel(input_shape=(32, 32, 1))
best_cnn.load_weights(cnn_path)
```

```
[ ]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at
0x7fc31608e908>
```

```
[ ]: # Predictions for best CNN Model
```

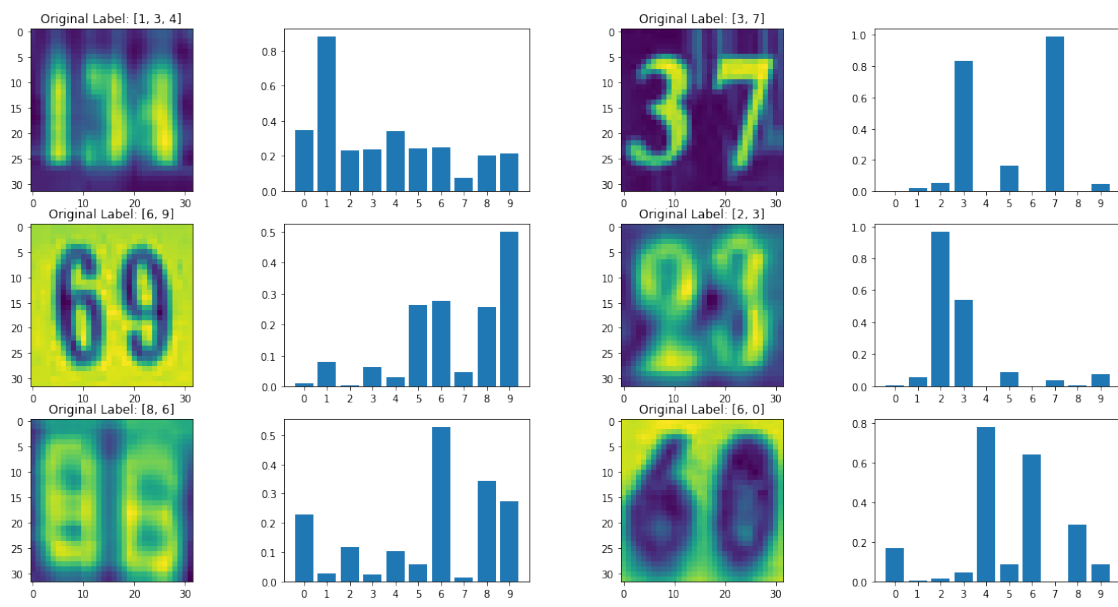
```
plt.rcParams["figure.figsize"] = 20, 10

idx = np.random.randint(0, SIZE["train"], 6)

for i in range(6):
    plt.subplot(3, 4, 2*i+1)
    plt.imshow(train_images[idx[i]])
    plt.title("Original Label: {}".format(train_labels[idx[i]]))

    plt.subplot(3, 4, 2*i + 2)
    plt.bar(x = np.arange(10), height=best_seq.predict(train_images[idx[i]] [np.
→newaxis, ..., np.newaxis])[0])
    plt.xticks(np.arange(10))

plt.show()
```



[: